

Exercise 4: Theoretical Problems

Student Number: 100269314

3. Triggers

- (a) Write CREATE TABLE commands for the relations. Add two checks: insuranceValue cannot be negative, and year must be between 1900–2100. None of the attributes can be NULL. Remember to define keys as well.

Solution.

```
CREATE TABLE Property(
    id INTEGER PRIMARY KEY,
    address VARCHAR(100) NOT NULL,
    city VARCHAR(100) NOT NULL,
    year INTEGER CHECK (year>=1900 AND year<=2100) NOT NULL,
    owner CHAR(11) REFERENCES Owner(ssNo) NOT NULL,
    insuranceValue FLOAT CHECK (insuranceValue>=0.0) NOT NULL
);

CREATE TABLE Owner(
    ssNo CHAR(11) PRIMARY KEY,
    name VARCHAR(100) NOT NULL,
    address VARCHAR(100) NOT NULL
);
```

- (b) Insert at least 10 rows for the tables, so triggers can be tested in part (f).

Solution.

	ssNo	name	address
1	12345678911	John	address 1
2	12345678912	John2	address 2
3	12345678913	John3	address 3
4	12345678914	John4	address 4
5	12345678915	John5	address 5
6	12345678916	John6	address 6
7	12345678917	John7	address 7
8	12345678918	John8	address 8
9	12345678919	John9	address 9
10	12345678910	John10	address 0

	id	address	city	year	owner	insuranceValue
1	1	address 1	CITY 1	2008	12345678910	19.2
2	2	address 2	CITY 1	2038	12345678911	190.2
3	3	address 3	CITY 1	1908	12345678911	100
4	4	address 4	CITY 2	1998	12345678910	200
5	5	address 5	CITY 1	2028	12345678915	300
6	6	address 6	CITY 1	2000	12345678913	400
7	7	address 7	CITY 4	2007	12345678910	500
8	8	address 8	CITY 1	2003	12345678912	77.77
9	9	address 9	CITY 5	2028	12345678915	123.45
10	10	address 10	CITY 6	2019	12345678910	19.3

- (c) When inserting a row into table Property, check that the city is written with all caps and fix if needed.

Solution.

```
CREATE TRIGGER CityAllCaps
    BEFORE INSERT
    ON Property
    FOR EACH ROW
    WHEN (NEW.city! = upper(NEW.city) )
BEGIN
```

```

INSERT INTO Property VALUES (
    NEW.id,
    NEW.address,
    upper(NEW.city),
    NEW.year,
    NEW.owner,
    NEW.insuranceValue
);
SELECT RAISE(IGNORE);
END;

--test 3c (part f - see images)
-- doesn't trigger CityAllCaps
INSERT INTO Property(id, address,city,year,owner,insuranceValue)
VALUES
(999, "address 1", "CITY 1", 2008, "12345678910", 19.2);
-- triggers CityAllCaps
INSERT INTO Property(id, address,city,year,owner,insuranceValue)
VALUES
(1000, "address 2", "ciTy 1", 2018, "12345678910", 19.1);

```

- (d) When updating a row in table Property, check that the new value of insuranceValue is at least 75% of the previous value. If it is smaller, set the new value of insuranceValue to 75% of the previous value.

Solution.

```

CREATE TRIGGER insuranceValueFloor
    AFTER UPDATE
    ON Property
    FOR EACH ROW
    WHEN (NEW.insuranceValue < OLD.insuranceValue * 0.75)
BEGIN
    UPDATE Property
    SET insuranceValue = OLD.insuranceValue * 0.75
    WHERE id = NEW.id;
END;

--test 3d (part f - see image)
-- doesn't trigger insuranceValueFloor
UPDATE Property SET insuranceValue = 70 WHERE id = 3;
-- triggers insuranceValueFloor
UPDATE Property SET insuranceValue = 175 WHERE id = 4;

```

- (e) When removing a row from table Owner, check that removed person is not owner of any property in the database. If person still owns something, prevent the removal.

Solution.

```

CREATE TRIGGER ownerDeleteProtection
    BEFORE DELETE

```

```

        ON Owner
    FOR EACH ROW
    WHEN EXISTS(
    SELECT 1
    FROM Owner, Property
    WHERE Property.owner=OLD.ssNo
    )
BEGIN
    SELECT RAISE(IGNORE);
END;
-- test 3e
-- doesn't trigger ownerDeleteProtection
DELETE FROM Owner WHERE ssNo=="12345678918";
-- triggers ownerDeleteProtection
DELETE FROM Owner WHERE ssNo=="12345678910";

```

- (f) After that, write two commands for each trigger (b), (c), and (d). The other should launch the trigger and the other should not. List the contents of the tables Property and Owner both before and after running all those statements (intermediate results are not required). For example, screenshot from SQLiteStudio is enough.

Solution.

3c - Before and After

	id	address	city	year	owner	insuranceValue
1	1	address 1	CITY 1	2008	12345678910	19.2
2	2	address 2	CITY 1	2038	12345678911	190.2
3	3	address 3	CITY 1	1908	12345678911	100
4	4	address 4	CITY 2	1998	12345678910	200
5	5	address 5	CITY 1	2028	12345678915	300
6	6	address 6	CITY 1	2000	12345678913	400
7	7	address 7	CITY 4	2007	12345678910	500
8	8	address 8	CITY 1	2003	12345678912	77.77
9	9	address 9	CITY 5	2028	12345678915	123.45
10	10	address 10	CITY 6	2019	12345678910	19.3

	id	address	city	year	owner	insuranceValue
1	1	address 1	CITY 1	2008	12345678910	19.2
2	2	address 2	CITY 1	2038	12345678911	190.2
3	3	address 3	CITY 1	1908	12345678911	100
4	4	address 4	CITY 2	1998	12345678910	200
5	5	address 5	CITY 1	2028	12345678915	300
6	6	address 6	CITY 1	2000	12345678913	400
7	7	address 7	CITY 4	2007	12345678910	500
8	8	address 8	CITY 1	2003	12345678912	77.77
9	9	address 9	CITY 5	2028	12345678915	123.45
10	10	address 10	CITY 6	2019	12345678910	19.3
11	999	address 1	CITY 1	2008	12345678910	19.2
12	1000	address 2	CITY 1	2018	12345678910	19.1

3d - Before and After

	id	address	city	year	owner	insuranceValue
1	1	address 1	CITY 1	2008	12345678910	19.2
2	2	address 2	CITY 1	2038	12345678911	190.2
3	3	address 3	CITY 1	1908	12345678911	100
4	4	address 4	CITY 2	1998	12345678910	200
5	5	address 5	CITY 1	2028	12345678915	300
6	6	address 6	CITY 1	2000	12345678913	400
7	7	address 7	CITY 4	2007	12345678910	500
8	8	address 8	CITY 1	2003	12345678912	77.77
9	9	address 9	CITY 5	2028	12345678915	123.45
10	10	address 10	CITY 6	2019	12345678910	19.3
11	999	address 1	CITY 1	2008	12345678910	19.2
12	1000	address 2	CITY 1	2018	12345678910	19.1

	id	address	city	year	owner	insuranceValue
1	1	address 1	CITY 1	2008	12345678910	19.2
2	2	address 2	CITY 1	2038	12345678911	190.2
3	3	address 3	CITY 1	1908	12345678911	75
4	4	address 4	CITY 2	1998	12345678910	175
5	5	address 5	CITY 1	2028	12345678915	300
6	6	address 6	CITY 1	2000	12345678913	400
7	7	address 7	CITY 4	2007	12345678910	500
8	8	address 8	CITY 1	2003	12345678912	77.77
9	9	address 9	CITY 5	2028	12345678915	123.45
10	10	address 10	CITY 6	2019	12345678910	19.3
11	999	address 1	CITY 1	2008	12345678910	19.2
12	1000	address 2	CITY 1	2018	12345678910	19.1

3e - Before and After

	ssNo	name	address
1	12345678911	John	address 1
2	12345678912	John2	address 2
3	12345678913	John3	address 3
4	12345678914	John4	address 4
5	12345678915	John5	address 5
6	12345678916	John6	address 6
7	12345678917	John7	address 7
8	12345678918	John8	address 8
9	12345678919	John9	address 9
10	12345678910	John10	address 0

	ssNo	name	address
1	12345678911	John	address 1
2	12345678912	John2	address 2
3	12345678913	John3	address 3
4	12345678914	John4	address 4
5	12345678915	John5	address 5
6	12345678916	John6	address 6
7	12345678917	John7	address 7
8	12345678919	John9	address 9
9	12345678910	John10	address 0

4. **Indices** Consider the relation schema OrderContent(orderID, product, amount) from the online store database introduced in the Exercise Round 1. Let's assume the relation occupies 120 pages of space. On average, each order includes 5 products and each product belongs to 30 orders. There is no clustering of any attributes.

Two kinds of queries are frequent for the table: searching for certain order ID (type Q1, fraction p1) and searching for orders with certain products (type Q2, fraction p2).

Insertions to the table take the fraction 1 - p1 - p2 of all operations on the table (type I).

Give formulas in terms of p1 and p2 to measure the cost of queries Q1 and Q2 and insertion I under the following four combinations (similarly to the exercise sessions).

- No indices at all

- Index for the attribute orderID
- Index for the attribute product
- Index for both attributes (orderID and product)

Solution.

No Index at All

$Q_1: 120$

⇒ Need to Scan whole table

$Q_2: 120$

⇒ Need to Scan whole table

$I: 2$

⇒ One access to read page & one to write it back modified

$$\text{Average: } 120p_1 + 120p_2 + (1-p_1-p_2)2 = 118p_1 + 118p_2 + 2$$

Index for attribute orderID

$Q_1: 6$

⇒ 1 to read index
+ avg S_{product} (i.e. p_k)
per order ID

$Q_2: 120$

⇒ Need to Scan whole table

$I: 4$

⇒ 2 to write new data +
2 to update index

$$\text{Average: } 6p_1 + 120p_2 + (1-p_1-p_2)4 = 2p_1 + 116p_2 + 4$$

Index for attribute product

$Q_1: 120$

⇒ Need to Scan whole table

$Q_2: 31$

⇒ 1 to read index
+ avg 30 orders
contain certain product

$I: 4$

⇒ 2 to write new data +
2 to update index

$$\text{Average: } 120p_1 + 31p_2 + (1-p_1-p_2)4 = 116p_1 + 27p_2 + 4$$

Index for both

$Q_1: 6$

⇒ 1 to read index
+ avg S_{product} (i.e. p_k)
per order ID

$Q_2: 31$

⇒ 1 to read index
+ avg 30 orders
contain certain product

$I: 6$

⇒ 2 to write new data +
4 to update indices


$$\text{Average: } 6p_1 + 31p_2 + (1-p_1-p_2)6 = 25p_2 + 6$$

5. Transactions

Solution.

a) i) '123' \Rightarrow stock = 65


Is possible : one such scenario (rep. as ordering of transactions)...



□

ii) '123' \Rightarrow stock < 50

Is possible : one such scenario (rep. as ordering of transactions)...



□


iii) table products contains only one tuple

Impossible : due to all transactions $T_1 \dots T_4$ uphold ACID principles, thanks to the atomicity of each transaction, T_4 will never run while T_3 is running (which is the only opportunity for deletion of tuples). Hence, the minimum # of tuples that can be in the table after the Transactions is 3 (w/ only mouse tuple removed)

□

iv) can order 'ABC shampoos' but not 'mouse'

Is possible : one such scenario (rep. as ordering of transactions)...




□


b) i) table products contains only one tuple w/o atomicity

Possible

one scenario



ii) Possible; one such scenario.



(Note: This depends on whether the deletion of T_3 is an atomic substep or also non-atomic. If it is non-atomic it is possible as indicated. BUT if the deletion occurs atomically, then it is not possible as mouse will always be removed & shampoo is orderable any time it is not removed as a whole)

iii) Yes possible. one scenario is shown in 5a(iii)