

CS-A1153 Databases: Vaccine Distribution Project

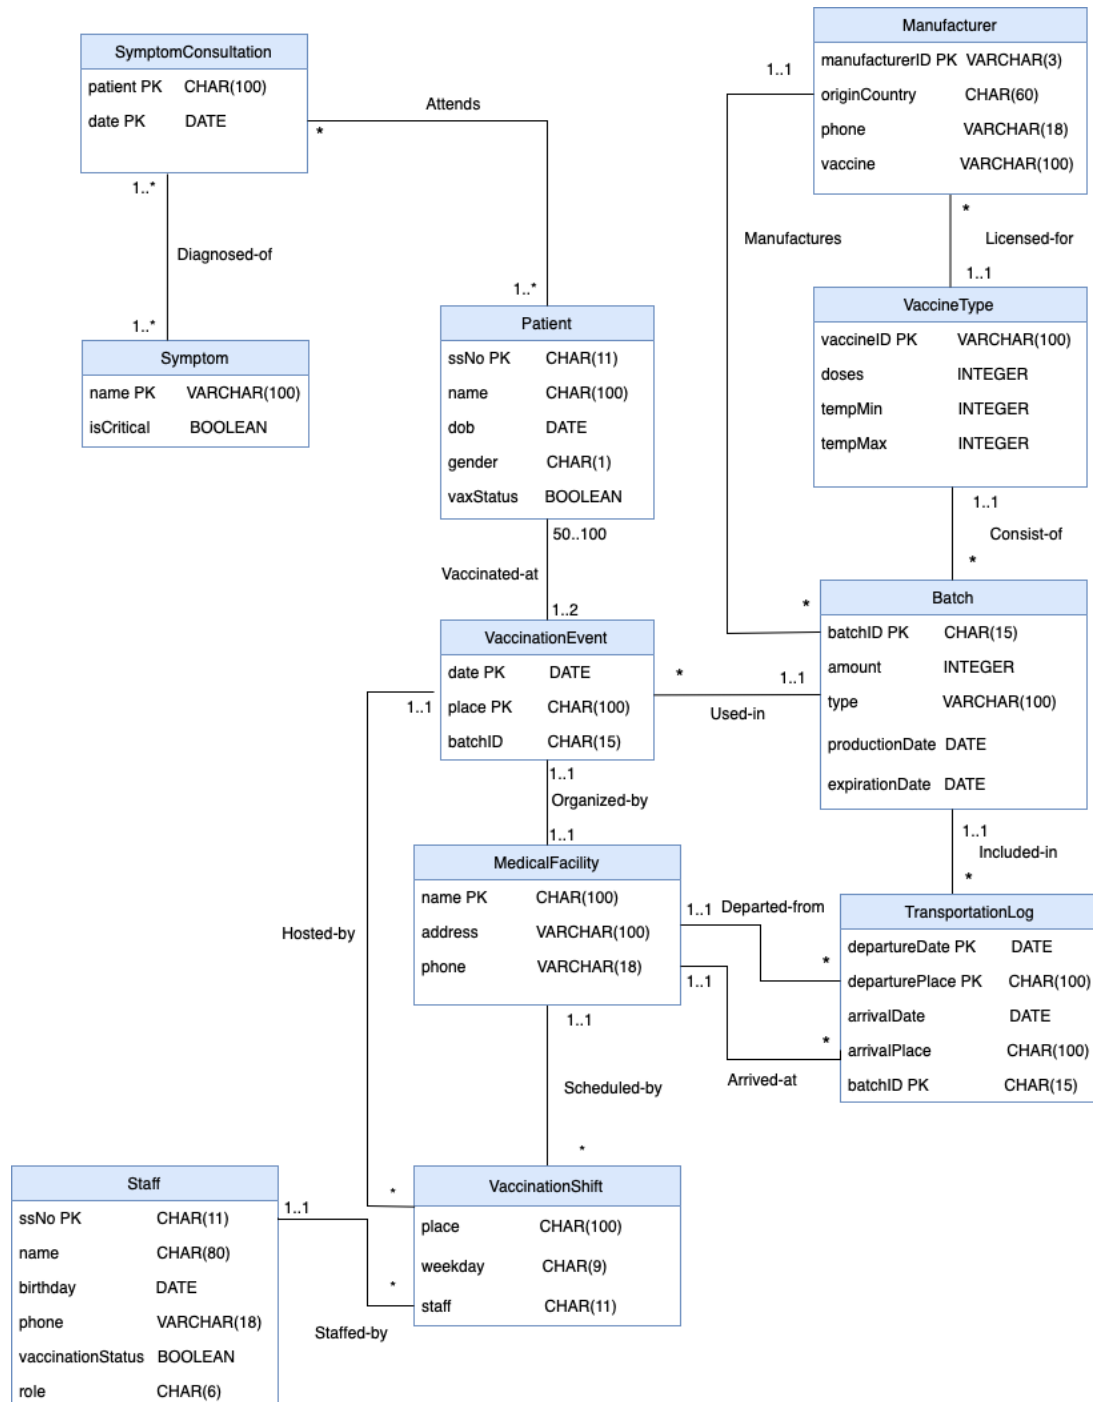
Group 9: Bin Choi, Karina Mina, Phuong Hoang, Tommaso Praturlon

Contents

1	Constructing the UML Diagram and Relational Data Model (Part I)	2
1.1	UML Diagram	2
1.1.1	Notes about Design Choice	3
1.2	Relational Data Model	5
1.3	Assessing the Model	5
1.3.1	Functional Dependencies	5
1.3.2	Redundancy and Anomaly	6
1.3.3	Boyce-Codd Normal Form	7
2	Modelling the Vaccine Distribution in Finland (Part II)	10
2.1	Database changes	10
2.2	Data cleaning	10
2.2.1	Empty columns and rows	10
2.2.2	Social security number in wrong format	11
2.2.3	Fix boolean type	11
2.2.4	Wrong dates format and ValueError: day out of month range	11
2.2.5	Extra columns in the data compared to our original UML design	12
2.3	SQL queries	13
2.3.1	Query #1	13
2.3.2	Query #2	13
2.3.3	Query #3	14
2.3.4	Query #4	15
2.3.5	Query #5	15
2.3.6	Query #6	17
2.3.7	Query #7	17
3	Appendix A	20

1 Constructing the UML Diagram and Relational Data Model (Part I)

1.1 UML Diagram



1.1.1 Notes about Design Choice

Manufacturer:

- The relation stores information about the manufacturer of vaccines.
- It is assumed that one manufacturer can produce only one vaccine type, and many manufacturers can produce the same vaccine type.

VaccineType:

- *vaccineID* is chosen as PK for **VaccineType** because it can be used to identify the vaccine type in associated tables.

Batch:

- This class stores the data about received vaccine batches: batch ID, number of vaccines, vaccine type, production and expiration date. *BatchID* is a PK because we assume it is the only unique attribute.

VaccinationEvent:

- This class stores the data about vaccination events: date, place and batch ID. According to the descriptions provided and our assumptions, on the same date there can be events in different places so date cannot be PK, at the same time same place can be used for events on different dates so place can't be PK either, same goes for the batchID, which can be used in several places and on several dates. However, on the same date in the same place there can only be one batchID, so *date* and *place* are both PKs.

Symptom:

- We first identify which attribute will be the PK for **Symptom** by asking some relevant questions:
 - Can one patient be diagnosed of the same symptom more than once (e.g. breathlessness initially diagnosed as not critical, next visit diagnosed as critical)?
 - Should we instead define the **symptom** relation as a **symptomConsultation** which is a subclass of a new class **Consultation** (which although is unused now, may increase scalability/expandability of this database in the future)?
 - Can we make the assumption that the symptom name locks in the criticality of the symptom?
- After inquiring with the TA and group mates, we realized that the last point is a relevant and fair assumption to make which can reduce the number of redundancies.
- Therefore, relation **Symptom** represents the symptom description/name and their criticality and is referenced by **SymptomConsultation**.

SymptomConsultation:

- This class will handle the data of patients' visit to hospital as result of experiencing symptoms.

- it stores date of diagnosis, name of symptoms, and the patient's ssNo.

Patient:

- There was not much difficulty with designing this relation as it is quite straightforward.
- This relation stores relevant particulars of patients who are receiving the vaccine.

TransportationLog:

- The table contains records about the movement of the batches between different medical facilities
- For this relation we assume that one batch does not move more than once a day. However, it may be possible that one batch is moved more than once in a week, therefore the selection of the composite PK.
- Another assumption made is that a tuple records only the initial and final place for the trip, not possible intermediate stops. For example, if the batch is moved from A to C through B, only A and C will appear in the tuple.
- One medical facility can move more than one batch on the same day.

MedicalFacility:

- the relation stores information about hospitals and clinics where the vaccine batches are stored and where vaccinations take place.

VaccinationShift:

- This table should contain information about the shifts that are organized at hospitals. Therefore it must be connected with both tables **MedicalFacility** and **Staff**. Because:
 - none of the fields **MedicalFacility** and **weekday** is unique, thereby leading this table to have anomalies, and
 - one staff, who can be a part-time worker, can work on different shifts at the same hospital,

it would be logical to include *hospital*, *weekday* and *staff* as fields of **VaccinationShift**. Field *staff* should be a foreign key connecting with table **Staff**'s primary key

Staff:

- As the table **staff** has contained *staff* that will be linking to this table, now we should decide the PK for **Staff**. The best candidate for PK is *ssNo* because it is unique to each personnel.

1.2 Relational Data Model

Manufacturer(manufacturerID, originCountry, phone, vaccineID)

VaccineType(vaccineID, doses, tempMin, tempMax)

Note: The assumption made for the association between **Manufacturer** and **VaccineType** is that many manufacturers can produce one type of vaccine.

Batch(batchID, amount, type, productionDate, expirationDate)

VaccinationEvent(date, place, batchID)

VaccinatedAt(date, place, ssNo)

Patient(ssNo, name, dob, gender, vaxStatus)

SymptomConsultation(patient, date, symptom)

Symptom(name, isCritical)

Note: The assumption made in the final relation **Symptom** is that the symptom name locks in the criticality of the symptom experienced. In other words, we assume that the symptom name (diagnosed and recorded by the doctor) is descriptive of the criticality.

MedicalFacility(name, address, phone)

VaccinationShift(place, weekday, staff)

Staff(ssNo, name, birthday, phone, vaccinationStatus, role)

TransportationLog(batchID, departureDate, departurePlace, arrivalDate, arrivalPlace)

1.3 Assessing the Model

1.3.1 Functional Dependencies

Note: In addition to the FD's explicitly listed below, the full list of FD's include all of the variations (of different combinations of dependants) that can be created by applying *Splitting & Combining FD rules* to the FD's below as explained during lecture.

Manufacturer

manufacturerID \rightarrow originCountry, phone, vaccine

phone \rightarrow manufacturerID, originCountry, vaccine

VaccineType

vaccineID \rightarrow doses, tempMin, tempMax

Batch

batchID \rightarrow amount, type, productionDate, expirationDate

VaccinationEvent

date, place \rightarrow batchID

VaccinatedAt

No FD as all attributes are PK.

Patient

ssNo \rightarrow name, dob, gender, vaxStatus

SymptomConsultation

No FD as all attributes are PK.

Symptom

name \rightarrow isCritical

MedicalFacility

name \rightarrow phone, address

phone \rightarrow name, address

TransportationLog

batchID, departureDate, departurePlace \rightarrow arrivalDate, arrivalPlace

batchID, arrivalPlace, arrivalDate \rightarrow departureDate, departurePlace

batchID, departureDate \rightarrow departurePlace, arrivalPlace, arrivalDate

batchID, arrivalDate \rightarrow departureDate, departurePlace, arrivalPlace

VaccinationShift

place \rightarrow staff

place, weekday \rightarrow staff

Staff

ssNo \rightarrow name, birthday, phone, vaccinationStatus, role

phone \rightarrow ssNo, name, birthday, vaccinationStatus, role

name, birthday \rightarrow ssNo, phone, vaccinationStatus, role

1.3.2 Redundancy and Anomaly**Manufacturer**

The relation has very little redundancy. The case in which the most tuples would be similar to each other would be when one manufacturer produces many different vaccines. However, given a realistic scenario in which only few vaccines are available, the redundancy would not cause problems.

VaccineType

There is little redundancy in this relation.

Batch

There is redundancy in this relation. Quite often the only variation between tuples would be the *batchID*.

VaccinationEvent

There is little to no redundancy in this relation (only date or place can sometimes appear several times).

VaccinatedAt

There is redundancy in this relation. Most often the only variation between tuples would be the *ssNo* because around 50-100 people are being vaccinated during one vaccination event (meaning same place and date attributes).

Patient

There is little to no redundancy in the relation Patient.

SymptomConsultation

There is some redundancy in this relation. For instance, if one patient is diagnosed of three different symptoms in one consultation, this will result in three different tuples with identical values in all attribute fields except for **symptom** field. This is a bearable amount of redundancy and there is not many ways to better represent the data.

Symptom

There is little to no redundancy in the relation Symptom.

TransportationLog

There is redundancy in the case in which one medical facility moves more than one batch on the same day to the same hospital or clinic. In this case the only variation between tuples would be the *batchID*. However, assuming that very few batches are moved at the same time from the same facility, the relation is acceptable.

MedicalFacility

In this relation there are not apparent anomalies or redundancies.

VaccinationShift

place and *weekday* are repeated because of *staff*. This is a redundancy that can cause deletion and update anomalies. However this is an acceptable redundancy because when we consider the functional dependencies (see **1.3.1**) and BCNF (see **1.3.3**), if we split this relation into smaller relations with available FDs, we will have two relations: one with *place* and *staff* and one with *place* and *weekday*. These relations will be too unnecessarily granular to portray vaccination shifts.

Staff

There is little redundancy in this relation.

1.3.3 Boyce-Codd Normal Form**Manufacturer:**

- **R(manufacturerID, originCountry, phone, vaccine):**
 - Considering the FD $\text{phone} \rightarrow \text{manufacturerID}, \text{originCountry}, \text{vaccine}$, here *phone* is a key therefore also a superkey
 - Therefore this relation is in BCNF

VaccineType:

- **R(vaccineID, doses, tempMin, tempMax):**
 - FD $\text{vaccineID} \rightarrow \text{doses}, \text{tempMin}, \text{tempMax}$ has *vaccineID* as a superkey
 - Therefore this relation is in BCNF

Batch:

- **R(batchID, amount, type, productionDate, expirationDate):**

- Consider FD $\text{batchID} \rightarrow \text{amount, type, productionDate, expirationDate}$. The LHS: *batchID* is a superkey of the relation
- Therefore this relation is in BCNF

VaccinationEvent:

- **R(date, place, batchID):**

- Consider FD $\text{date, place} \rightarrow \text{batchID}$. The LHS: *date, place* is a superkey of the relation
- Therefore this relation is in BCNF

VaccinatedAt:

- **R(date, place, ssNo):**

- There are no FD's for this relation.
- Therefore this relation is in BCNF

Patient:

- **R(ssNo, name, dob, gender):**

- Consider FD $\text{ssNo} \rightarrow \text{name, dob, gender, vaxStatus}$. The LHS: *ssNo* is a superkey of the relation
- Therefore this relation is in BCNF

Symptom:

- **R(name, isCritical):**

- Consider FD $\text{name} \rightarrow \text{isCritical}$. The LHS: *name* is a superkey of the relation
- Therefore this relation is in BCNF

SymptomConsultation:

- **R(patient, date, symptom):**

- There are no FD's for this relation.
- Therefore this relation is in BCNF

MedicalFacility:

- **R(name, phone, address):**

- Considering the FD $\text{phone} \rightarrow \text{name, address}$ *phone* is a key therefore also a superkey
- Therefore this relation is in BCNF

TransportationLog:

- **R(batchID, departureDate, departurePlace, arrivalDate, arrivalPlace):**
 - Considering the FD $\text{batchID} \rightarrow \text{arrivalPlace}, \text{arrivalDate}$, the left-hand side of the dependency is a superkey
 - the same is valid of all the other nontrivial FDs
 - Therefore this relation is in BCNF

VaccinationShift:

- **R(place, weekday, staff):**
 - FD $\text{place} \rightarrow \text{staff}$ holds in this relation but the left side is not a superkey.
 - Therefore this relation is not in BCNF.
- **R1(place, staff):**
 - FD $\text{place} \rightarrow \text{staff}$ has place as a superkey in this relations
 - Therefore this relation is in BCNF.
- **R2(place, weekday):**
 - There is no FD in this relation
 - Therefore this relation is in BCNF.

Staff:

- **R(ssNo, name, birthday, phone, vaccinationStatus, role):**
 - FD $\text{ssNo} \rightarrow \text{name, birthday, phone, vaccinationStatus, role}$ has ssNo as a superkey
 - FD $\text{phone} \rightarrow \text{ssNo, name, birthday, vaccinationStatus, role}$ has phone as a superkey
 - FD $\text{name, birthday} \rightarrow \text{ssNo, phone, vaccinationStatus, role}$ has name, birthday as a composite superkey
 - Therefore this relation is in BCNF

2 Modelling the Vaccine Distribution in Finland (Part II)

2.1 Database changes

Overall, we updated many data types in Part II for the purpose of consistency. All attributes with variable length became `varchar(100)` e.g name, location. All attributes with fixed number of characters became `char(*)`, where `*` is the number of characters e.g `ssNo` is always `char(11)`.

Staff

According to the feedback from previous assignment, *Staff* has been added another column of *place* where the staff works. A foreign key was also added to connect relation *Staff* with *MedicalFacility*

Batch

A new one-to-many relation was added **is-stored-in** between *Batch* and *MedicalFacility* with `*` on *Batch*, so location was added as an attribute and FK to relation *Batch*. We also noticed that relation *Batch* was missing manufacturer from *Manufacturer* among its attributes although there was a one-to-many association between these two relations already in our UML diagram in Part I, so we added it along with location when creating tables in database.

2.2 Data cleaning

The data cleaning process can be described shortly as:

- First, we iterate through the whole Excel and return a dictionary that contains key as sheet name and value as sheet content. Each of these values is a pandas dataframe.
- Second, we iterate through this dictionary's values and clean them using functions that we have made.

The whole process's codes in detail can be found in file `process_raw.py`

Below are the data errors found while inserting into relations the content provided and what we have done to handle the errors.

2.2.1 Empty columns and rows

There were some empty columns and rows from the Excel files. This was handled by `.dropna()` function applied to both axes. Just in case we also trim spaces from values.

```
def trim_df(df=None):  
    if df is None:  
        print("Dataframe must be passed")  
    else:  
        df.dropna(inplace=True, how='all', axis='columns')  
        df.dropna(inplace=True, how='all')  
        df.applymap(lambda x: x.strip() if isinstance(x, str) else x)
```

2.2.2 Social security number in wrong format

The social security numbers listed are not in the correct format of how Finnish social security number should be. Strings were simply split at their positions and pasted together to return the correct format.

```
def process_ssNo(df):
    for col in df.columns.values:
        if check_is_ssNo(col) is True:
            df[['first', 'second']] = df[col].str.split('-', expand=True)

            if len(df[col][0]) == 13:
                df['year'] = df['first'].str[2:4]
                df['month'] = df['first'].str[3:5]
                df['day'] = df['first'].str[5:7]
            else:
                df['year'] = df['first'].str[:2]
                df['month'] = df['first'].str[2:4]
                df['day'] = df['first'].str[4:6]

            df[col] = df['day'] + df['month'] + df['year'] + '-' + df['second']

    df = df.drop(['first', 'second', 'day', 'month',
                 'year'], axis=1, inplace=True)
```

check_is_ssNo() is a helper function to determine whether the column contains social security numbers.

2.2.3 Fix boolean type

The format of boolean in Excel file was 1s and 0s. We check if columns only has 1s and 0s and turn them into type bool

```
for col in df.columns.values:
    if df[col].isin([1,0]).all():
        df[col] = df[col].astype(bool)
```

2.2.4 Wrong dates format and ValueError: day out of month range

We noticed there are some date errors in sheet *Diagnosis*. To handle the error, we wrote function handle_dates and apply to all columns that we check could be dates. The idea of this function is:

- First, we iterate through the rows in the column. If the row is integer then we update row with date. Formula is date 20 December 1899 + value of the row.
- Second, we check if the date is actually date by trying to extract from row value year value. Doing this instead of the normal checking of the full date format is because we simply do not care what exactly is the format of dates. There are two reasons why we do not care:

- Even if pandas could not pass the data as `dtype_datetime64` any due to some errors, it always recognizes date format from Excel and pass in as `YYYY-MM-DD` and maybe something else after that object.
- Luckily, postgres is also generally easy with date type and will pass in any date format as date, as long as it is not a timestamp, which does not appear in our case.
- If the extraction return error, we catch this error and handle by simply splitting the object at day position (because object is always in format `YYYY-MM-DD`) and subtract by 1. Then we update the row with the full string we just fixed.

```
def handle_dates(df, col_name):
    for row in df[col_name]:
        if isinstance(row, numbers.Number):
            df[col_name].mask(df[col_name] == row, datetime.date(1899, 12, 30) +
                               datetime.timedelta(days=row) , inplace=True)
        else:
            try:
                row.strftime('%Y')
            except Exception:
                df[col_name].mask(df[col_name] == row, str(row)[:8] +
                                   str(int(str(row)[8:10]) - 1) + str(row)[10:], inplace=True)
```

2.2.5 Extra columns in the data compared to our original UML design

This is not a data error per se but is included because it interferes with inserting the data into sql relation smoothly and we cannot edit the raw data to fit our UML. This is handled by getting relations' number of columns from `information_schema` and bind the same amount of variables and dataframe columns in our insert sql query. This step was included as a class `Postgres` method instead of in `process_raw.py` because it has to do more on Postgres side.

```
def execute_insert(self, df=None, sql_table=None):
    if df is None or sql_table is None:
        print("All parameters must be passed.")
    else:
        try:
            # get amount of columns that we define in our tables
            sql_table = sql_table.lower()
            self.execute_single_sql('select count(*) as column_count from
                                     information_schema.columns where table_schema = \'public\'
                                     and table_name = %s', (sql_table,))
            sql_cols = self.cursor.fetchone()[0]

            # insert only the first amount of columns that we define
            sql = "insert into " + sql_table + " values(" +
                  bind_string(int(sql_cols)) + ")"

            # truncate table
            trunc_sql = 'truncate table ' + sql_table + ' cascade'
```

```

self.execute_single_sql(trunc_sql)

df = df[df.columns[:sql_cols]]

for i in range(len(df.index)):
    self.execute_single_sql(sql, df.values.tolist()[i],
                           commit=True)

except (Exception, Error) as e:
    raise

```

2.3 SQL queries

2.3.1 Query #1

This query is to find all staff members who work at vaccination on May 10, 2021. This was achieved by joining relations *Staff*, *VaccinationEvent* and *VaccinationShift*. Because there is only *date* in *VaccinationEvent*, *weekday* and *date* must be turned into numbers to join.

```

1 SELECT Staff.* FROM Staff JOIN
2     (SELECT staff, place,
3      CASE WHEN weekday = 'Sunday' THEN 0
4      WHEN weekday = 'Monday' THEN 1
5      WHEN weekday = 'Tuesday' THEN 2
6      WHEN weekday = 'Wednesday' THEN 3
7      WHEN weekday = 'Thursday' THEN 4
8      WHEN weekday='Friday' THEN 5
9      WHEN weekday = 'Saturday' THEN 6 END AS dow_no
10     FROM VaccinationShift) shift
11 ON Staff.ssNo = shift.staff
12 JOIN (SELECT date, extract(dow from date) AS wd, place FROM
13      vaccinationevent) event
14 ON shift.dow_no = event.wd and shift.place = event.place
15 WHERE event.date = '2021-05-10';

```

The reference results for the latter query are reported in Table 1.

2.3.2 Query #2

The query below is to find the doctors that are available on Wednesdays in Helsinki. The author interpreted the request as finding the doctorS who will be working in vaccination event on Wednesday shifts at a hospitals in Helsinki. To achieve this the relations *Staff*, *VaccinationShift* and *MedicalFacility* are joined together.

```

1 SELECT Staff.name FROM Staff JOIN
2 VaccinationShift ON Staff.ssNo = VaccinationShift.staff
3 JOIN MedicalFacility ON VaccinationShift.place = MedicalFacility.name

```

```

4 WHERE VaccinationShift.weekday = 'Wednesday' AND Staff.role = 'doctor'
   AND MedicalFacility.address LIKE '%HELSINKI%';

```

The reference results for the latter query are reported in Table 2.

2.3.3 Query #3

The intent of this query was to state the current location and the last location of a vaccine batch and compare them for consistency check. The solution applied involved a subquery which returned for each *batchID* in the relation *TransportationLog* the date of the last movement of the batch, along with its ID. This result is later intersected with the *TransportationLog* table for the ID and date and from this query the last location is taken, while from the relation *Batch* the current location and ID are kept.

```

1 SELECT T1.batchID,
2     location AS currentLocation,
3     T1.arrivalPlace AS lastLocation
4 FROM Batch,
5     TransportationLog AS T1
6     INNER JOIN (
7         SELECT batchID, MAX(arrivalDate) AS lastDate
8         FROM TransportationLog
9         GROUP BY batchID
10    ) AS T2
11    ON T1.batchID = T2.batchID
12     AND T1.arrivalDate = lastDate
13 WHERE Batch.batchID = T1.batchID
14 ;

```

The reference results for the latter query are reported in Table 3.

The second part of the query asked to list the *batchID* where the current location was not consistent with the report in the transportation log, and the phone number of the facility where it was supposed to be found. For this reason the following query makes use of the previous one (3a) and compares the *currentLocation* and *lastLocation*. When these are different the query returns the desired output. The results can be referenced in Table 4.

```

1 SELECT Q1.batchID, phone
2 FROM MedicalFacility,
3     (SELECT T1.batchID,
4         location AS currentLocation,
5         T1.arrivalPlace AS lastLocation
6     FROM Batch,
7         TransportationLog AS T1
8         INNER JOIN (
9             SELECT batchID,
10                 MAX(arrivalDate) AS lastDate
11             FROM TransportationLog
12             GROUP BY batchID
13         ) AS T2

```

```

14     ON T1.batchID = T2.batchID
15     AND T1.arrivalDate = lastDate
16     WHERE Batch.batchID = T1.batchID
17 ) AS Q1
18 WHERE Q1.currentLocation != Q1.lastLocation
19     AND Q1.lastLocation = MedicalFacility.name;

```

2.3.4 Query #4

With this query the intent is to find all patients with critical symptoms diagnosed later than May 10, 2021 and match this data with the data about the vaccines the patient has been given (such as batches of the vaccines, the type of the vaccine, the date the vaccine was given, and the location of the vaccination). Two subqueries were used to retrieve the needed information. From the relation *Symptom* only critical symptoms were selected, then all the information regarding the patient has been selected from relations *Patient* and *SymptomConsultation* where also the filtering for date and symptom was applied. Lastly, the outer query selects patient's name and ssNo from the results generated by inner queries as well as it also selects the vaccination related data such as batchID, vaccine type, vaccination date and location from relations *VaccinationEvent*, *VaccinatedAt* and *Batch*. The query produces 0 results, which are displayed in Table 5. The output should be correct because there were no critical symptoms diagnosed after May 10, 2021. If we use the same query and change the date filtering for earlier than May 10, 2021, then it generates results properly, which means that the query itself works as intended and produces desired results.

```

1  SELECT P.ssNo, P.name, VaccinationEvent.batchID, Batch.type,
   VaccinatedAt.date, VaccinatedAt.place
2  FROM (SELECT Patient.name, Patient.ssNo, SymptomConsultation.symptom
3         FROM Patient, SymptomConsultation
4         WHERE Patient.ssNo = SymptomConsultation.patient AND
5              SymptomConsultation.date > '2021-05-10'
6         GROUP BY name, ssNo, symptom
7         HAVING SymptomConsultation.symptom IN(
8              SELECT name FROM Symptom WHERE isCritical is True)
9  ) AS P, VaccinationEvent, Batch, VaccinatedAt
10 WHERE VaccinatedAt.ssNo=P.ssNo AND VaccinationEvent.batchID =
   Batch.batchID AND
10 VaccinationEvent.place=VaccinatedAt.place AND
   VaccinationEvent.date=VaccinatedAt.date
11 ;

```

2.3.5 Query #5

With this query the intent is to create a view for patients with additional column "vaccinationStatus". This column takes the value 1 if the patient has attended enough vaccinations, and 0 otherwise. Although all of the vaccine types in the data provided require two doses, I do not think it is appropriate to check that each patient received two doses

(irregardless of the vaccine type) as new vaccine types could be added to the database in the future. Then, the query will be rendered useless. Therefore, I designed my query such that it considers the vaccine type that the patient has received and checks whether they received the necessary number of dosages. The results are displayed in Table 6.

```

1 CREATE VIEW patientVaxStatus AS (
2 SELECT
3     ssno,
4     name,
5     dob,
6     gender,
7     CASE
8         WHEN vaxcount = doses THEN 1
9         WHEN vaxcount != doses THEN 0
10    END vaxStatus
11 FROM
12     patient
13 NATURAL JOIN (
14     SELECT
15         *
16     FROM
17         (
18         SELECT
19             ssno,
20             COUNT(*) AS vaxCount
21         FROM
22             vaccinatedat v
23         GROUP BY
24             ssno) AS t1
25 NATURAL JOIN (
26     SELECT
27         DISTINCT ssno,
28         doses
29     FROM
30         (
31         SELECT
32             ssno,
33             type
34         FROM
35             (vaccinatedat
36             NATURAL JOIN vaccinationevent) AS v,
37             batch b
38         WHERE
39             v.batchid = b.batchid) AS t1
40     INNER JOIN vaccinetype ON
41         vaccineid = type) AS t2)
42 AS T2
43 );
44 -- to see all tuples of the view...
```

```
45 SELECT * FROM patientVaxStatus;
```

2.3.6 Query #6

With this query the intent is to find the total number of vaccines stored in each medical facility and the number of distinct vaccine types available. From the relation *Batch* it is possible to retrieve all these information with the following query, grouping by medical facility. The results are displayed in Table 7.

```
1 SELECT location,
2     SUM(amount) AS totVax,
3     COUNT(DISTINCT type) AS vaxTypes
4 FROM Batch
5 GROUP BY location
6 ;
```

2.3.7 Query #7

With this query the intent is to find the average frequency of different symptoms diagnosed for each vaccine type/name. The symptom should not be considered to be caused by the vaccine, if it has been diagnosed before the patient got the vaccine. If a patient has received two different types of vaccines before the diagnosis of the symptom, the symptom should be counted once for both of the vaccines. The key to coming up for this query was creating the intermediate output T1 and picking out the relevant statistics from there. Also, we were advised to make the assumption that all occurrences of a symptom report is counted for more accurate result (e.g. same symptom diagnosis of one patient is counted twice, one patient can be diagnosed of multiple symptoms, etc...). The results are displayed in Table 8.

```
1 WITH
2 T1 AS (
3 SELECT
4     v.ssno AS pid,
5     name AS VaccineName,
6     symptom
7 FROM
8     (
9     vaccinatedat v
10 INNER JOIN VaccinationEvent e ON
11     e.place = v.place
12     and e.date = v.date
13 NATURAL JOIN Batch
14 INNER JOIN vaccinetype ON
15     vaccinetype.vaccineid = Batch.type
16     ),
17     symptomconsultation c
18 WHERE
19     c.patient = v.ssno
```

```

20     AND c.date > v.date
21 ),
22 T2 AS (
23 SELECT
24     VaccineName,
25     COUNT(*) AS count
26 FROM
27     T1
28 GROUP BY
29     VaccineName
30 ),
31 T3 AS (
32 SELECT
33     VaccineName,
34     symptom,
35     COUNT(*) AS count
36 FROM
37     T1
38 GROUP BY
39     VaccineName,
40     symptom
41 )
42 SELECT
43     T2.VaccineName,
44     T3.symptom,
45     ROUND(1.0 * T3.count / T2.count, 4) AS Frequency
46 FROM
47     T3
48 INNER JOIN T2 ON
49     T3.VaccineName = T2.VaccineName;

```

Note: to identify by vaccine type (id) rather than vaccine name, we simply replace VaccineName with VaxType. The results are displayed in Table 9.

```

1  with
2  T1 as (
3  select
4      v.ssno as pid,
5      vaccinetype.vaccineid as VaxType,
6      symptom
7  from
8      (
9      vaccinatedat v
10 inner join VaccinationEvent e on
11     e.place = v.place
12     and e.date = v.date
13 natural join Batch
14 inner join vaccinetype on
15     vaccinetype.vaccineid = Batch.type

```

```
16     ),
17     symptomconsultation c
18 where
19     c.patient = v.ssno
20     and c.date > v.date
21 ),
22 T2 as (
23 select
24     VaxType,
25     COUNT(*) as count
26 from
27     T1
28 group by
29     VaxType
30 ),
31 T3 as (
32 select
33     VaxType,
34     symptom,
35     COUNT(*) as count
36 from
37     T1
38 group by
39     VaxType,
40     symptom
41 )
42 select
43     T2.VaxType,
44     T3.symptom,
45     ROUND(1.0 * T3.count / T2.count, 4) as Frequency
46 from
47     T3
48 inner join T2 on
49     T3.VaxType = T2.VaxType;
```

3 Appendix A

<i>ssno</i>	<i>name</i>	<i>birthday</i>	<i>phone</i>	<i>role</i>	<i>vaccination status</i>	<i>place</i>
802092-4854	Kaden Tromp	1992-08-02	044-624-1591	nurse	t	Tapiola Health Center
914074-7140	Deon Hoppe	1974-09-19	040-399-1121	nurse	f	Tapiola Health Center
614094-4448	Jordy Hilpert	1994-06-15	044-506-1982	doctor	t	Tapiola Health Center
813063-6581	Jazlyn Schneider	1963-08-12	040-868-2528	nurse	t	Sanomala Vaccina- tion Point
007177-5988	Samir Hills	1977-10-03	040-093-0059	nurse	t	Sanomala Vaccina- tion Point
818088-8027	Haylie Wintheiser	1988-08-17	050-448-8894	nurse	t	Myyrmki Energia Areena
212082-5928	Elena Bartell	1982-02-18	041-938-9451	nurse	t	Myyrmki Energia Areena
222072-1761	Alfreda Champlin	1972-02-23	041-631-1851	nurse	t	Myyrmki Energia Areena

Table 1: Query 1

<i>name</i>
Rosalia Simonis
Shaylee Kris
Hilbert Purdy
Elnora Greenholt

Table 2: Query 2

<i>batchID</i>	<i>currentLocation</i>	<i>lastLocation</i>
B01	Sanomala Vaccination Point	Sanomala Vaccination Point
B02	Messukeskus	Sanomala Vaccination Point
B03	Myyrmäki Energia Areena	Myyrmäki Energia Areena
B04	Malmi	Malmi
B06	Iso Omena Vaccination Point	Myyrmäki Energia Areena
B07	Myyrmäki Energia Areena	Myyrmäki Energia Areena
B08	Tapiola Health Center	Tapiola Health Center
B12	Sanomala Vaccination Point	Sanomala Vaccination Point
B13	Iso Omena Vaccination Point	Iso Omena Vaccination Point
B15	Malmi	Malmi
B16	Tapiola Health Center	Tapiola Health Center
B17	Myyrmäki Energia Areena	Myyrmäki Energia Areena
B18	Tapiola Health Center	Tapiola Health Center
B21	Iso Omena Vaccination Point	Iso Omena Vaccination Point
B22	Myyrmäki Energia Areena	Myyrmäki Energia Areena
B23	Sanomala Vaccination Point	Sanomala Vaccination Point
B24	Malmi	Malmi
B25	Malmi	Malmi
B27	Myyrmäki Energia Areena	Myyrmäki Energia Areena
B28	Iso Omena Vaccination Point	Iso Omena Vaccination Point
B29	Myyrmäki Energia Areena	Sanomala Vaccination Point
B30	Iso Omena Vaccination Point	Iso Omena Vaccination Point

Table 3: Query 3a

<i>batchID</i>	<i>phone</i>
B02	093-105-3153
B06	093-104-5930
B29	093-105-3153

Table 4: Query 3b

<i>ssno</i>	<i>name</i>	<i>batchid</i>	<i>type</i>	<i>date</i>	<i>place</i>
(0 rows)					

Table 5: Query 4

<i>ssno</i>	<i>name</i>	<i>dob</i>	<i>gender</i>	<i>varstatus</i>
291284-112N	Rodolfo O'Reilly	1984-12-29	M	1
140278-1893	Prof. Erling Morar MD	1978-02-14	F	0
030395-191X	Dr. Simeon Keeling II	1995-03-03	M	0
180273-253D	Dereck Beer	1973-02-18	M	0
141297-2818	Prof. Brice Metz PhD	1997-12-14	M	0
250306-323X	Darlene Brakus	2006-03-25	F	0
140699-395X	Josefa Greenfelder DVM	1999-06-14	M	0
180205-4796	Ms. Hassie Runolfsson PhD	2005-02-18	F	0
270100-4899	Ms. Opal Lang	2000-01-27	F	0
301072-5216	Noah Leuschke	1972-10-30	M	0

040189-753F	Lukas Runolfsson V	1989-01-04	M	1
010872-8748	Braeden Hackett	1972-08-01	M	0
011002-957O	John Larkin	2002-10-01	M	0
080679-9686	Gisselle Hilpert	1979-06-08	F	0
050884-1135	Lonzo Collier	1984-08-05	M	1
041174-114G	Marvin Fahey	1974-11-04	M	0
221173-126T	Johanna McClure	1973-11-22	F	0
270906-1438	Zetta Runolfsson	2006-09-27	F	0
150385-155F	Andreanne Jakubowski	1985-03-15	M	0
010615-1657	Mrs. Ophelia Corwin Sr.	2015-06-01	F	0
020209-1778	Alvera Medhurst	2009-02-02	F	0
060193-189U	Julius Marks	1993-01-06	M	0
171170-199K	Leland Moen	1970-11-17	M	0
070314-203V	Rhea Hettinger	2014-03-07	M	0
271103-2165	Mathew Buckridge	2003-11-27	M	0
010897-218B	Jonathan Wyman	1997-08-01	M	0
060788-240U	Prof. Willard Marquardt II	1988-07-06	M	0
181187-242U	Dr. Victor Armstrong	1987-11-18	M	0
290911-252V	Lilly Farrell V	2011-09-29	F	0
301004-267L	Isabell Nader	2004-10-30	F	0
140272-2797	Abbey Schuppe	1972-02-14	F	0
170107-2839	Estrella Johns	2007-01-17	F	0
111275-287B	Taylor Krajcik	1975-12-11	F	1
290706-292C	Marilou Ryan	2006-07-29	F	0
070709-295R	Alysson Jakubowski	2009-07-07	F	0
101090-3010	Emerald Johnson	1990-10-10	F	0
210406-302M	Faustino Barton	2006-04-21	M	1
081180-303B	Reva Waelchi	1980-11-08	F	0
010107-326Q	Orpha Bogisich	2007-01-01	F	0
140508-3385	Dakota Greenfelder	2008-05-14	F	0
100888-358W	Braxton Hane	1988-08-10	M	1
210672-378H	Ms. Alisha Ortiz	1972-06-21	F	0
030579-394M	Kathlyn Moore	1979-05-03	F	0
080593-413K	Mr. Reid Little II	1993-05-08	M	0
290696-4156	Rossie Spinka	1996-06-29	F	0
040619-440B	Prof. Kevon Cummings	2019-06-04	M	0
160409-443L	Aliyah Harber	2009-04-16	M	1
050213-474D	Elenora Sawayn	2013-02-05	F	0
040893-509I	Fay Ryan	1993-08-04	F	0
281187-519R	Flossie Torp	1987-11-28	F	1
220699-5231	Sid Hahn	1999-06-22	M	0
270301-525G	Dr. Mireille Hansen	2001-03-27	M	1
271170-5340	Mrs. Lorena Kreiger	1970-11-27	F	0
121018-5367	Bulah Heidenreich	2018-10-12	M	0
240771-5480	Dashawn Schamberger	1971-07-24	M	0
051081-5518	Miss Charity Powlowski	1981-10-05	F	0
260209-5673	Dr. Lamont Ferry	2009-02-26	M	1
200883-576C	Loyal Hoeger	1983-08-20	M	0

180312-5791	Mathilde Smith	2012-03-18	F	0
010201-5814	Harrison Heaney	2001-02-01	M	0
081204-5838	Dr. Margarette Mertz IV	2004-12-08	F	0
300916-586P	Aiden Volkman	2016-09-30	F	1
020916-592P	Britney Gutmann	2016-09-02	F	0
031101-6045	Mrs. Kailyn Collier DVM	2001-11-03	F	0
191188-6103	Destiny Konopelski PhD	1988-11-19	F	0
070689-6113	Prof. Raphael Prosacco DVM	1989-06-07	M	0
070896-613R	Omer Denesik	1996-08-07	M	0
131282-6162	Tabitha Howe	1982-12-13	M	0
250300-6271	Mariam Ritchie	2000-03-25	F	0
221104-6308	Prof. Demarco Hahn	2004-11-22	M	0
150619-6325	Ms. Nelda Brekke PhD	2019-06-15	F	0
040618-6419	Dr. Freddie Cartwright	2018-06-04	M	0
060500-642P	Cassandra Mayert	2000-05-06	F	0
301102-649D	Dr. Jayson Glover DVM	2002-11-30	M	0
030999-6514	Eldred Blanda	1999-09-03	F	0
061088-6543	Trycia Jaskolski	1988-10-06	F	0
060406-686D	Dovie West	2006-04-06	F	0
200411-6983	Nellie Nitzsche	2011-04-20	F	0
040893-7021	Prof. Harrison Toy	1993-08-04	M	0
080776-708H	Maia Towne II	1976-07-08	M	0
281285-732X	Shanna Osinski	1985-12-28	F	0
180321-737O	Corine Hane	2021-03-18	F	1
100385-787I	Cristal Borer	1985-03-10	M	0
020609-7898	Hillard Boehm V	2009-06-02	M	0
131104-8113	Prof. Yessenia Dooley Jr.	2004-11-13	F	0
180509-869W	Spencer Kunde	2009-05-18	M	0
221274-8947	Devon Nicolas	1974-12-22	M	0

Table 6: Query 5: created view is shown using **select** query (after view creation)

<i>location</i>	<i>totVax</i>	<i>vaxTypes</i>
Iso Omena Vaccination Point	65	3
Malmi	65	3
Messukeskus	120	3
Myyrmäki Energia Areena	85	3
Sanomala Vaccination Point	40	2
Tapiola Health Center	55	2

Table 7: Query 6

<i>vaccinename</i>	<i>symptom</i>	<i>frequency</i>
Comirnaty	high fever	0.0500
AstraZeneca	itchiness near injection	0.0750
AstraZeneca	muscle ache	0.2000
AstraZeneca	joint pain	0.1500
AstraZeneca	warmth near injection	0.0750
Moderna	feelings of illness	0.1667
Comirnaty	fever	0.1500
Comirnaty	diarrhea	0.1500
Comirnaty	pain near injection	0.0500
Moderna	chills	0.0417
AstraZeneca	inflammation near injection	0.0250
AstraZeneca	diarrhea	0.0250
Moderna	lymfadenopathy	0.0833
Moderna	headache	0.0417
Moderna	nausea	0.0833
Comirnaty	muscle ache	0.1500
AstraZeneca	fatigue	0.0250
Moderna	fever	0.0833
Moderna	high fever	0.0417
Comirnaty	joint pain	0.1000
AstraZeneca	feelings of illness	0.0250
AstraZeneca	headache	0.1500
Moderna	fatigue	0.0417
AstraZeneca	blurring of vision	0.0250
Moderna	joint pain	0.1667
Comirnaty	inflammation near injection	0.0500
AstraZeneca	nausea	0.0750
AstraZeneca	vomiting	0.0250
AstraZeneca	high fever	0.0500
AstraZeneca	fever	0.0750
Comirnaty	chest pain	0.0500
Comirnaty	fatigue	0.0500
Moderna	muscle ache	0.2083
Moderna	vomiting	0.0417
Comirnaty	headache	0.2000

Table 8: Query 7

<i>vaxtype</i>	<i>symptom</i>	<i>frequency</i>
V03	joint pain	0.1000
V03	chest pain	0.0500
V01	vomiting	0.0250
V01	warmth near injection	0.0750
V02	chills	0.0417
V01	headache	0.1500
V03	high fever	0.0500
V01	diarrhea	0.0250
V02	fever	0.0833
V02	high fever	0.0417
V02	nausea	0.0833
V02	muscle ache	0.2083
V01	high fever	0.0500
V03	fever	0.1500
V02	headache	0.0417
V02	joint pain	0.1667
V01	blurring of vision	0.0250
V01	muscle ache	0.2000
V03	headache	0.2000
V01	itchiness near injection	0.0750
V02	vomiting	0.0417
V01	feelings of illness	0.0250
V01	joint pain	0.1500
V03	muscle ache	0.1500
V03	inflammation near injection	0.0500
V01	fever	0.0750
V03	pain near injection	0.0500
V03	diarrhea	0.1500
V02	fatigue	0.0417
V01	inflammation near injection	0.0250
V02	feelings of illness	0.1667
V01	fatigue	0.0250
V03	fatigue	0.0500
V02	lymfadenopathy	0.0833
V01	nausea	0.0750

Table 9: Query 7b: Using a slightly different SQL query, we can compute query 7 to identify by vaccine type